

Every edit stops the playtest

THE USER

Technical designers at small studios iterate on playable prototypes with real playtesters in the room.

THE COST

Every stop costs the thing they are trying to observe: how people actually play.

Beetle's promise


Change the game without stopping the game.

TODAY'S LOOP



Each turn of the loop loses the live playtest you were trying to watch.

The Beetle workflow



WHAT IT RUNS ON

One machine

a single GB10 on the studio LAN

One local model

Ollama on loopback

One real agent runtime

OpenClaw with Beetle tools

Zero cloud

no fallback exists

The working video



[VIDEO PLACEHOLDER]

90 to 120 second demo recording, dropped in after rehearsal
Any time compression is labelled on screen

ON-SCREEN CAPTIONS

0 to 6 s Beetle. Change the game without stopping the game.

6 to 22 s A local model composes the world. Validated by game code before it is published.

22 to 34 s Two players join from their phones. Same screen, their phones only send controls.

34 to 58 s Edits land live. No disconnect, no reset.


58 to 90 s The validator refuses a world nobody can finish. The agent repairs it and commits only a passing version.

90 to 105 s Measured on one GB10. Local model, local agent, no cloud.

105 to 115 s Beetle

Fresh world, two phones, a live edit, preserved progress, a conflicting edit caught by the validator, the repaired version committed.

Agent architecture and validation



TWO KINDS OF STATE

- WorldSpec is structure. It changes only when a transaction commits, one version per commit.
- SessionState is players, relics, score. The simulation advances it; the agent never writes it.
- Commit needs a server-issued proof.

VALIDATION IS GAME CODE

- Geometry, bridge sockets, walk field, reachability with the gate locked, occupied support at commit time.
- An LLM saying valid is never approval.
- Nothing from the model runs. Output is schema-checked JSON only.

Measured local-first results

MODEL

**[model tag],
[quantization]**

on the GB10 via Ollama on loopback

BRIEF TO PLAYABLE

[min] / [p50] / [max] s

over [N] runs, cold load [cold s] reported separately

EDIT TO COMMIT

[min] / [p50] / [max] s

edit request to committed version

INVALID EDITS CAUGHT

[caught] of [attempts]

repairs that succeeded: [repaired]

SESSION CONTINUITY

**[players kept] players,
[relics kept] relics**

[reconnects] reconnects across [N commits] commits

OFFLINE PROOF

[what was disconnected]

still worked: [what still worked]

Every bracketed value is a placeholder. Values come only from data/benchmarks/.json and data/reports/*.json. Nothing unmeasured goes on a slide.*

Business hypothesis and next validation



HYPOTHESIS

- Studios will pay for faster prototype iteration that keeps playtests running.
- Their data stays on their machines, which matters to studios with unreleased IP.

NOT CLAIMED

- Market size.
- Pricing.
- Customer traction.
- Speed-up percentages without a measured baseline.

NEXT VALIDATION

- Time a conventional edit-rebuild-regather loop with the same designers and the same change list.
- Compare it against Beetle's measured edit-to-commit times.



Change the game without stopping the game.